


Federated training lands close to pooling the data; DP-SGD buys privacy for a little accuracy



DP-SGD run at final-round epsilon = 1.44 (delta = 1e-5) per operator.